

```
#ARIMA/SARIMA
```

```
library(plyr)
library(fpp3)
```

```
## -- Attaching packages ----- fpp3 0.4.0 --
```

```
## v tibble      3.1.2      v tsibble      1.0.1
## v dplyr       1.0.7      v tsibbledata 0.3.0
## v tidyr       1.1.3      v feasts       0.2.2
## v lubridate   1.7.10     v fable        0.3.1
## v ggplot2     3.3.5
```

```
## -- Conflicts ----- fpp3_conflicts --
```

```
## x dplyr::arrange()      masks plyr::arrange()
## x dplyr::count()       masks plyr::count()
## x lubridate::date()     masks base::date()
## x dplyr::failwith()     masks plyr::failwith()
## x dplyr::filter()       masks stats::filter()
## x dplyr::id()           masks plyr::id()
## x tsibble::intersect() masks base::intersect()
## x tsibble::interval()  masks lubridate::interval()
## x dplyr::lag()          masks stats::lag()
## x dplyr::mutate()       masks plyr::mutate()
## x dplyr::rename()       masks plyr::rename()
## x tsibble::setdiff()    masks base::setdiff()
## x dplyr::summarise()    masks plyr::summarise()
## x dplyr::summarize()    masks plyr::summarize()
## x tsibble::union()      masks base::union()
```

```
library(tsibble)
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

```
library(zoo)
```

```
##
## Attaching package: 'zoo'
```

```
## The following object is masked from 'package:tsibble':
##
##   index
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```

#read in the interpolated data
data <- readr::read_csv(file = 'data/data_interpolated_with_lags.csv') %>%
  mutate(yw = yearweek(yw)) %>%
  select(-X1) %>%
  as_tsibble(key = c(Mode, ORegionDAT, DRegionDAT), index = yw)

```

```
## Warning: Missing column names filled in: 'X1' [1]
```

```

##
## -- Column specification -----
## cols(
##   .default = col_double(),
##   yw = col_character(),
##   Mode = col_character(),
##   ORegionDAT = col_character(),
##   DRegionDAT = col_character()
## )
## i Use 'spec()' for the full column specifications.

```

```

#make into univariate approx_cost series

```

```

data <- data %>%
  select(Mode, ORegionDAT, DRegionDAT, yw, approx_cost, tmax_lag_12, tmax_lag_2, prcp_lag_12, prcp_lag_2)
  filter(Mode == "R", DRegionDAT == "IL_CHI")

```

```

#trim leading and trailing na's

```

```

data <- drop_na(data)

```

```

#create training set - up through 2020 of the time series

```

```

train <- data %>%
  filter_index(~ "2018 W52")

```

There are many possible arima models - based on choice of hyperparameters and whether to include seasonality or not. The ARIMA() function automatically chooses the best hyperparameters.

```

fit <- train %>%
  model(ARIMA(approx_cost ~ tmax_lag_12 + tmax_lag_2 + prcp_lag_12 + prcp_lag_2 + diesel_price + new_deaths))

```

```

## Warning: Provided exogenous regressors are rank deficient, removing regressors:
## 'new_deaths', 'pandemic'

```

```

#see what the automatically chosen arima models were.
report(fit)

```

```

## Series: approx_cost
## Model: LM w/ ARIMA(2,0,0) errors
##
## Coefficients:
##          ar1          ar2 tmax_lag_12 tmax_lag_2 prcp_lag_12 prcp_lag_2
##          1.0859 -0.2193          0.0008          0.0018         -0.1615         -0.0764
## s.e.    0.1030   0.1044          0.0055          0.0033          1.2465          0.2509

```


```
##      diesel_price  volume
##           0.5823    6e-04
## s.e.         0.1149    1e-03
##
## sigma^2 estimated as 0.007878:  log likelihood=95.61
## AIC=-173.22   AICc=-171.03   BIC=-150.53
```

It looks like for the first 3 time series ARIMA() automatically picked up on the period 52 seasonality. For Boston V data, it did not. Also it chose a different set of hyperparameters for each time series.

```
#forecast
#in order to produce the forecast of approx_cost, we need to feed in a forecast of tmax, prcp, and dies
future_data <- data %>%
  filter_index("2019 W01"~"2019 W26") %>%
  select(Mode, ORegionDAT, DRegionDAT, yw, approx_cost, tmax_lag_12, tmax_lag_2, prcp_lag_12, prcp_lag_2)
fc <- fit %>% forecast(future_data)
```

```
#plot
fc %>%
  autoplot(train) +
  autolayer(data, colour = "black") +
  labs(title="SARIMA")
```


```
## Plot variable not specified, automatically selected '.vars = approx_cost'
```



```
#Looking at fitted values and residuals
```

```
#get fitted values and residuals
aug = augment(fit)
```


```
#autoplot them
autoplot(aug, .innov) +
  labs(title = "Residuals from the ARIMA/SARIMA model")
```



```
#histograms
aug %>%
  ggplot(aes(x = .innov)) +
  geom_histogram() +
  facet_grid(rows = vars(Mode), cols = vars(DRegionDAT)) +
  labs(title = "Histograms of residuals")
```


```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Histograms of residuals



```
#acf
aug %>%
  ACF(.innov, lag_max = Inf) %>%
  autoplot() +
  facet_grid(rows = vars(Mode), cols = vars(DRegionDAT)) +
  labs(title = "ACF of Residuals")
```

ACF of Residuals



Looking at forecast errors

```
accuracy(fc, data)
```

```
## # A tibble: 1 x 13
##   .model Mode ORegionDAT DRegionDAT .type    ME  RMSE  MAE  MPE  MAPE  MASE
##   <chr>  <chr> <chr>      <chr>      <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 "ARIMA~ R    CA_FRS    IL_CHI    Test  -0.210 0.254 0.212 -12.9  13.0 0.586
## # ... with 2 more variables: RMSSE <dbl>, ACF1 <dbl>
```